



Service Fabric - Use Case Scenarios and Examples

Microsoft Services



Conditions and Terms of Use

Microsoft Confidential

This training package is proprietary and confidential, and is intended only for uses described in the training materials. Content and software is provided to you under a Non-Disclosure Agreement and cannot be distributed. Copying or disclosing all or any portion of the content and/or software included in such packages is strictly prohibited.

The contents of this package are for informational and training purposes only and are provided "as is" without warranty of any kind, whether express or implied, including but not limited to the implied warranties of merchantability, fitness for a particular purpose, and non-infringement.

Training package content, including URLs and other Internet website references, is subject to change without notice. Because Microsoft must respond to changing market conditions, the content should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information presented after the date of publication. Unless otherwise noted, the companies, organizations, products, domain names, e-mail addresses, logos, people, places, and events depicted herein are fictitious, and no association with any real company, organization, product, domain name, e-mail address, logo, person, place, or event is intended or should be inferred.

Copyright and Trademarks

© 2016 Microsoft Corporation. All rights reserved.

Microsoft may have patents, patent applications, trademarks, copyrights, or other intellectual property rights covering subject matter in this document. Except as expressly provided in written license agreement from Microsoft, the furnishing of this document does not give you any license to these patents, trademarks, copyrights, or other intellectual property.

Complying with all applicable copyright laws is the responsibility of the user. Without limiting the rights under copyright, no part of this document may be reproduced, stored in or introduced into a retrieval system, or transmitted in any form or by any means (electronic, mechanical, photocopying, recording, or otherwise), or for any purpose, without the express written permission of Microsoft Corporation.

For more information, see **Use of Microsoft Copyrighted Content** at
<https://www.microsoft.com/en-us/legal/intellectualproperty/permissions/default.aspx>

Microsoft®, Internet Explorer®, Outlook®, SkyDrive®, Windows Vista®, Zune®, Xbox 360®, DirectX®, Windows Server® and Windows® are either registered trademarks or trademarks of Microsoft Corporation in the United States and/or other countries. Other Microsoft products mentioned herein may be either registered trademarks or trademarks of Microsoft Corporation in the United States and/or other countries. All other trademarks are property of their respective owners.

How to View This Presentation

- All slides have two styles of PDF documents:
 - Full View – full page view per slide
 - Notes View – 3 slides per page with ruled lines used to take notes
- Use Adobe Reader or other PDF viewing tools to review the PDF slides
- All PDF slides are located in your AzureSFWS.zip file



Service Fabric - Use Case Scenarios and Examples

Applying Microservices on the Microsoft Platform

Microsoft Services



300+ Service Fabric Customers



Agenda

- TalkTalk
- Mesh Systems
- Illyriad Games



Service Fabric - Use Case Scenarios and Examples

TalkTalk TV

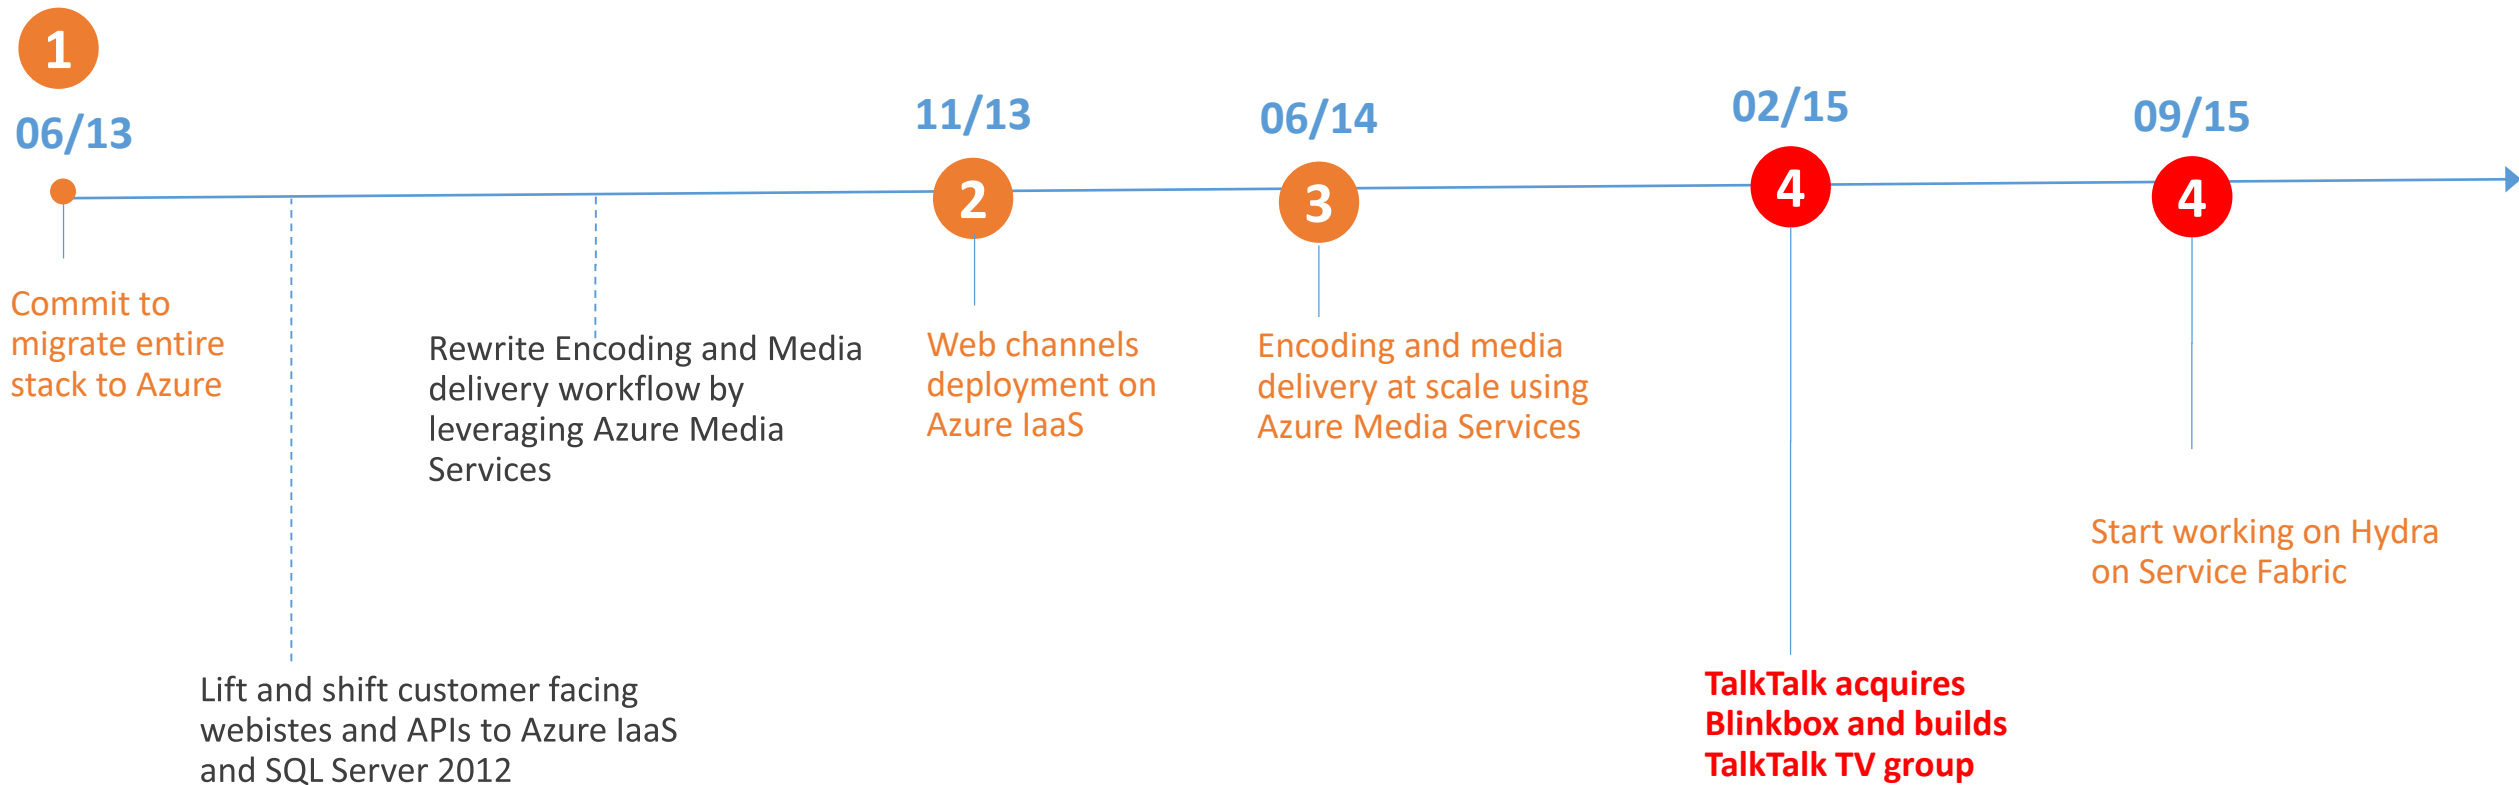
Microsoft Services





- UK's third largest cable TV Provider
- The Leading value for money phone, broadband and TV Provider for UK homes.
- With more than 1.4M customers, TalkTalk is the #3 pay TV platform in the UK
- TalkTalk TV delivers the latest TV and movie content to TalkTalk subscribers, Apple computers, PCs, game consoles, tablets, and smart TVs.

TalkTalk TV – A long Journey to the Cloud



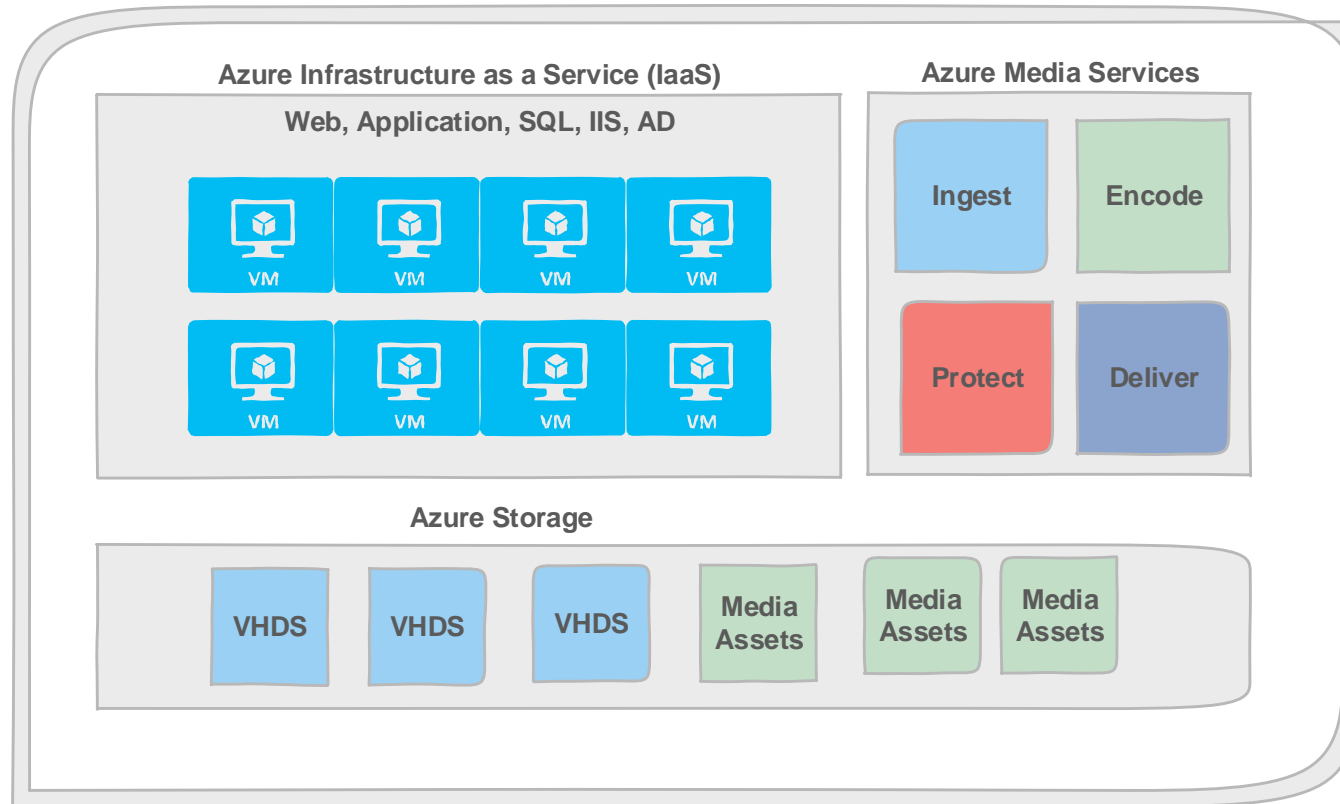
Hydra

Single Encoding Pipeline

batly – Lead Engineer, TalkTalk TV

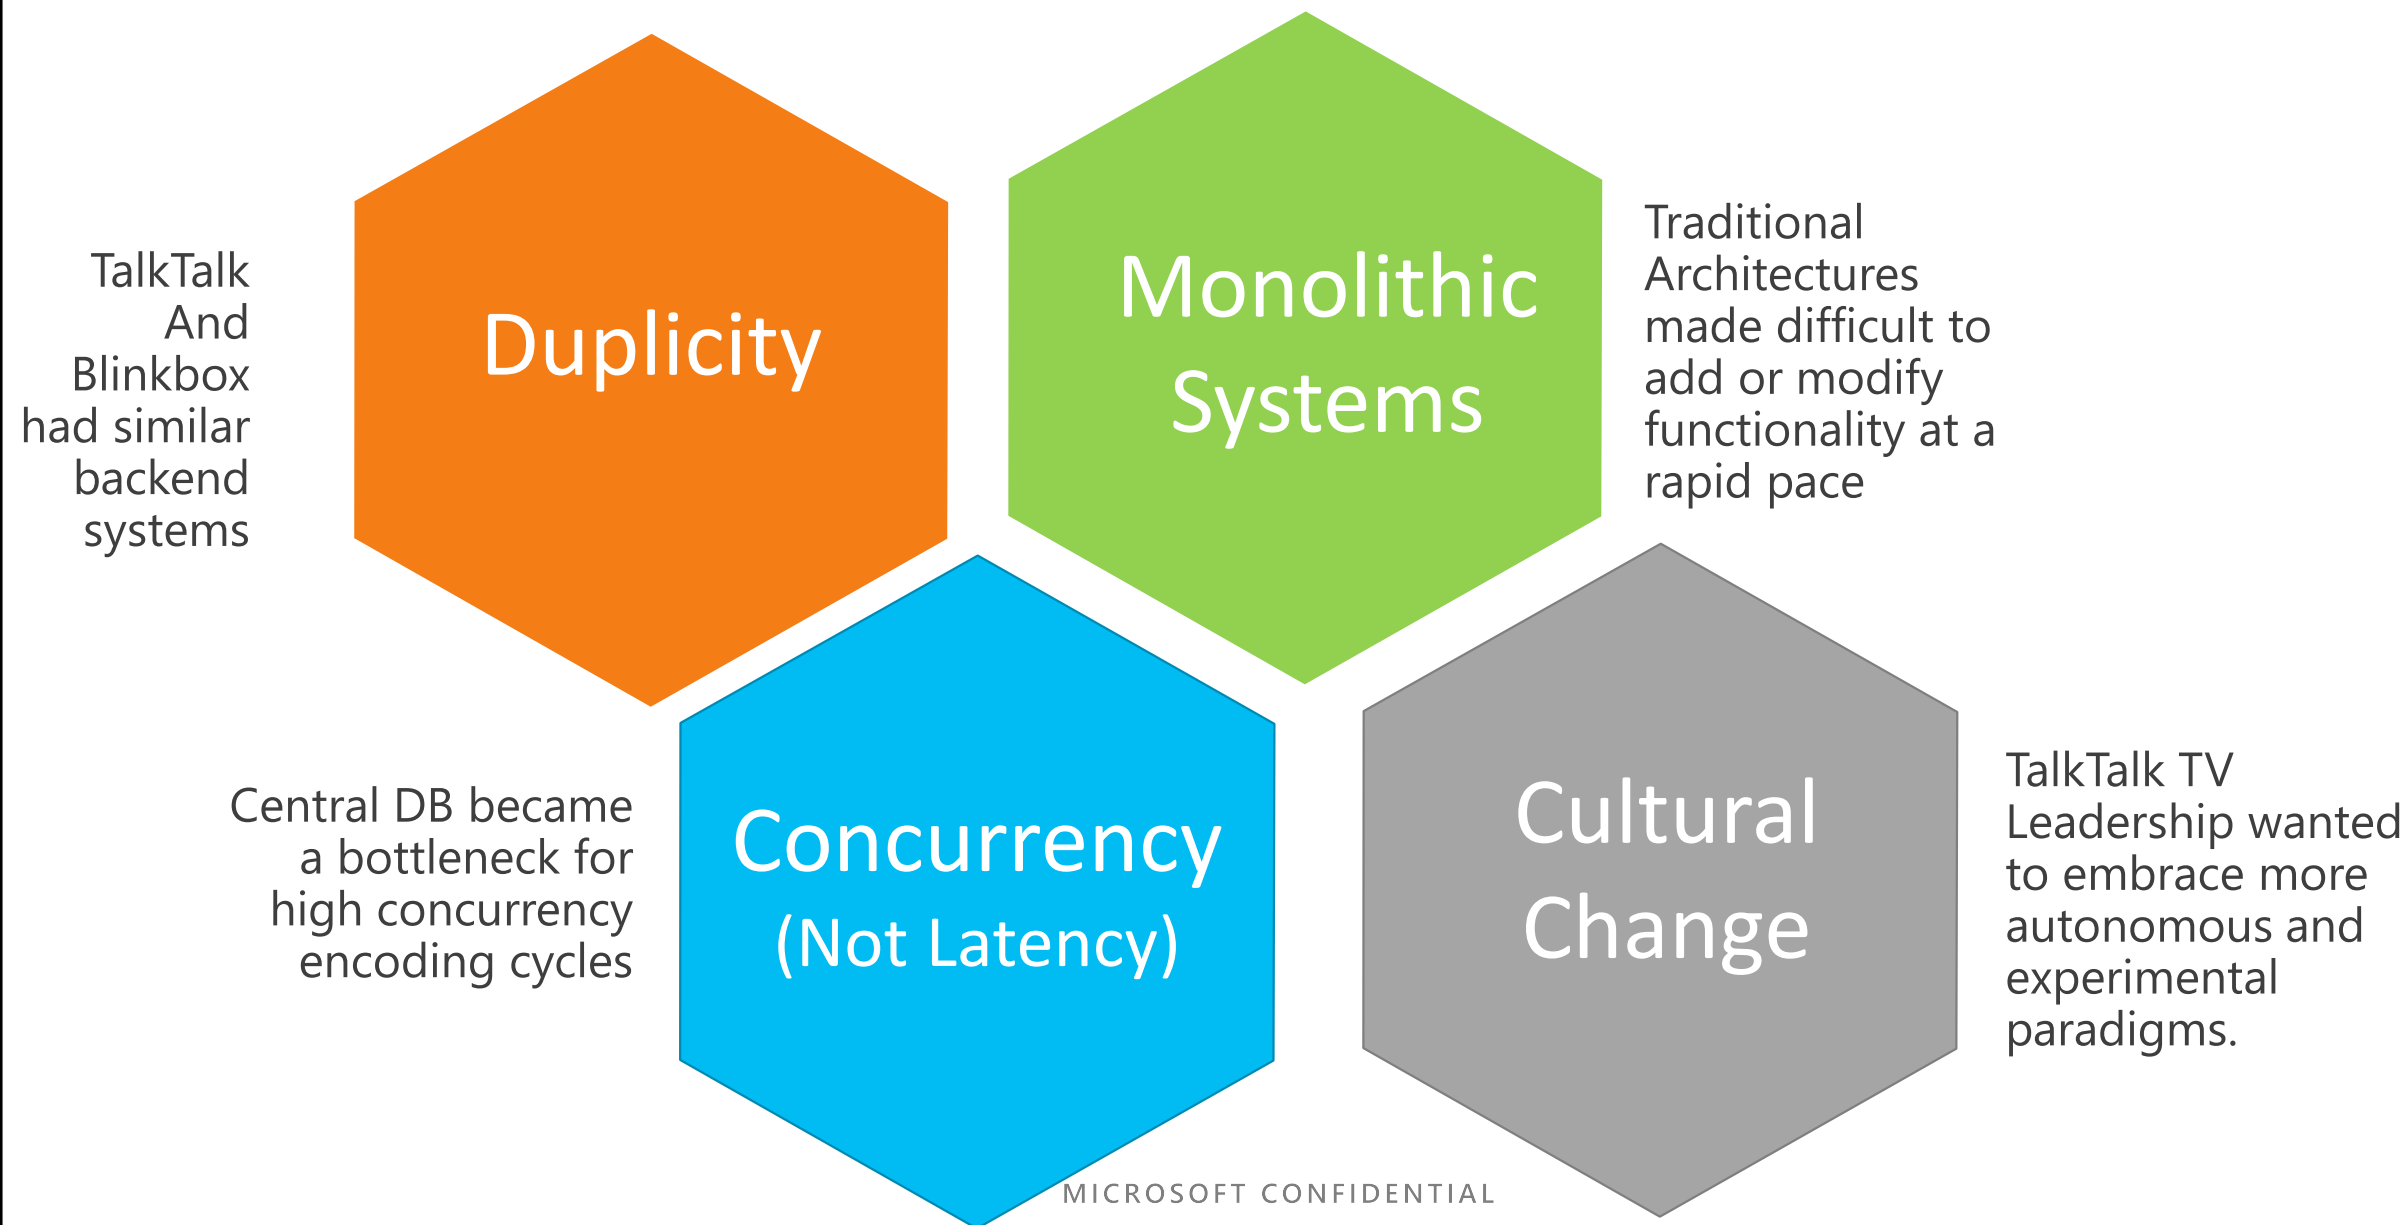


TalkTalk TV As IaaS

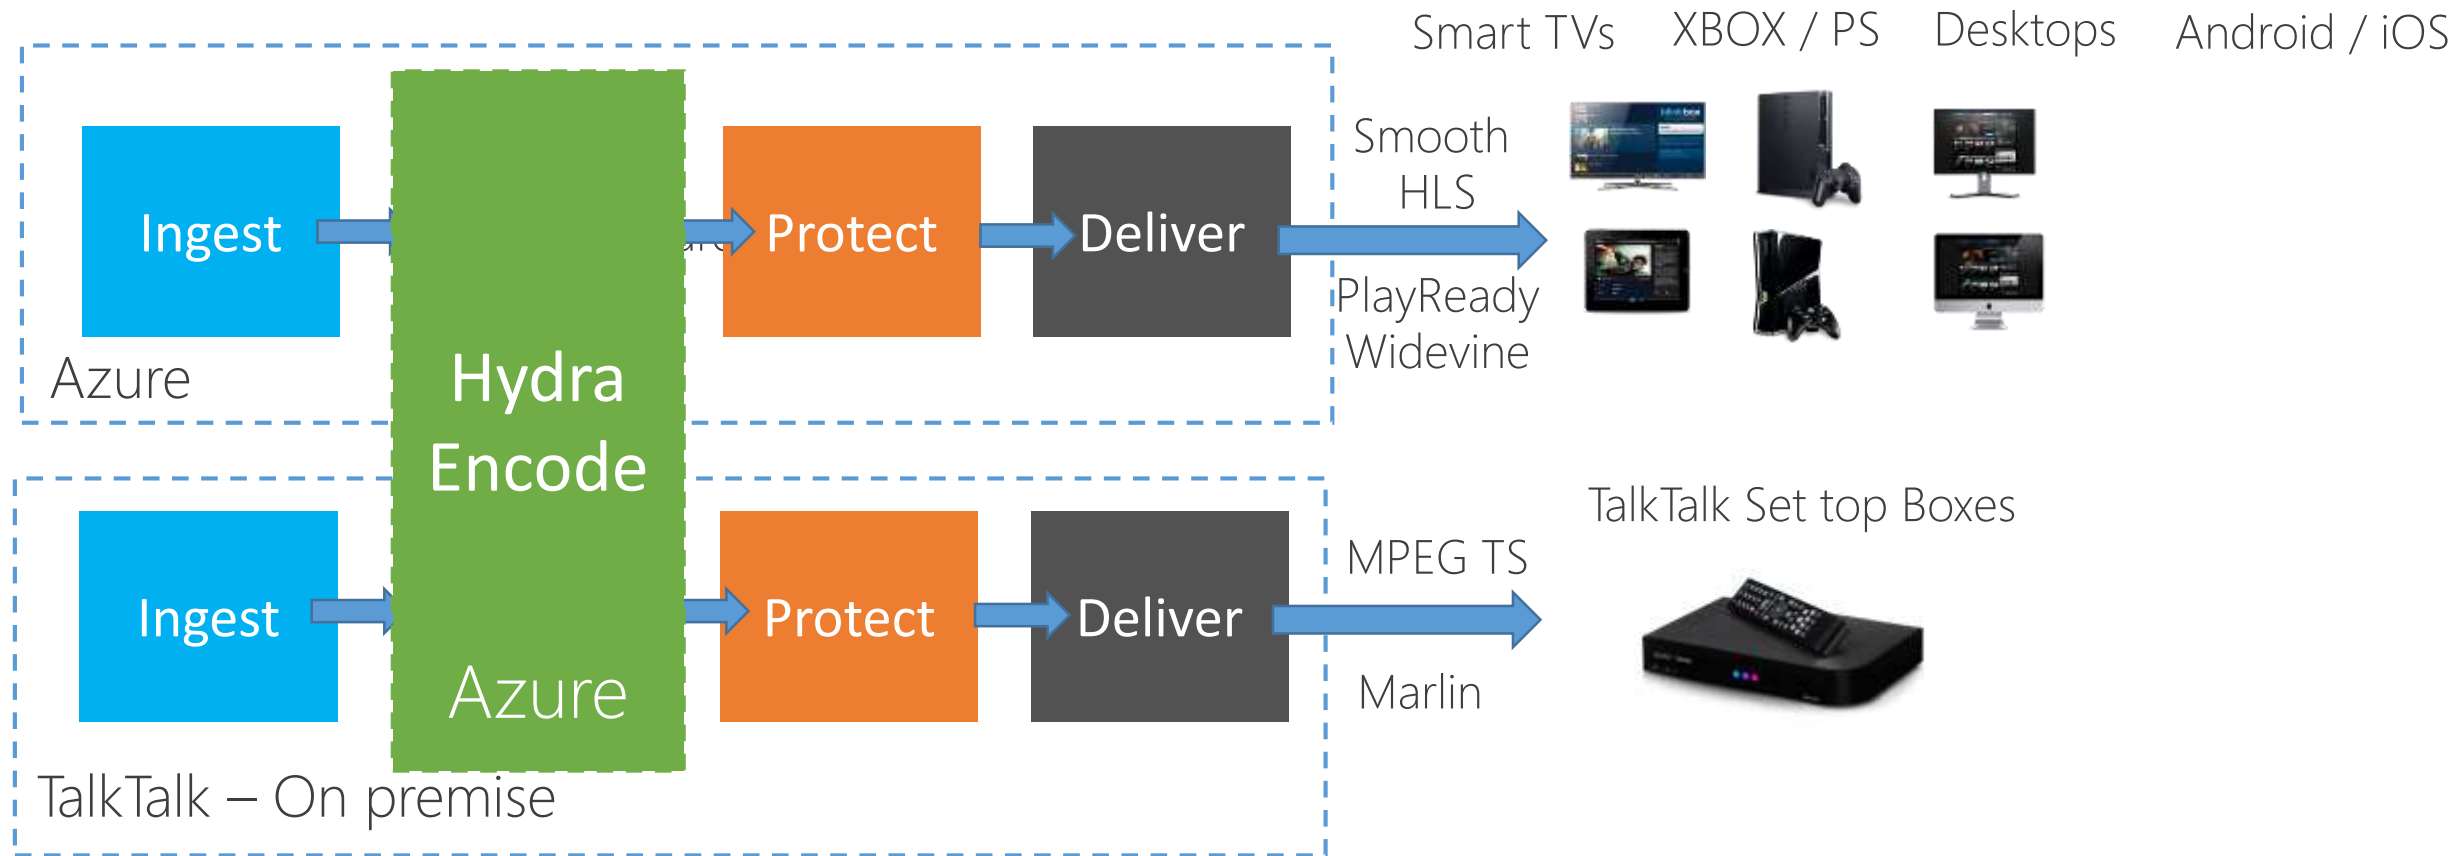


- > 200 virtual machine cores
- > 1.5 PB of Azure Storage
- > Media assets delivered from Azure NEU and WEU
- > Hybrid SQL Cluster (2 copies in the Cloud and an additional on-premises)

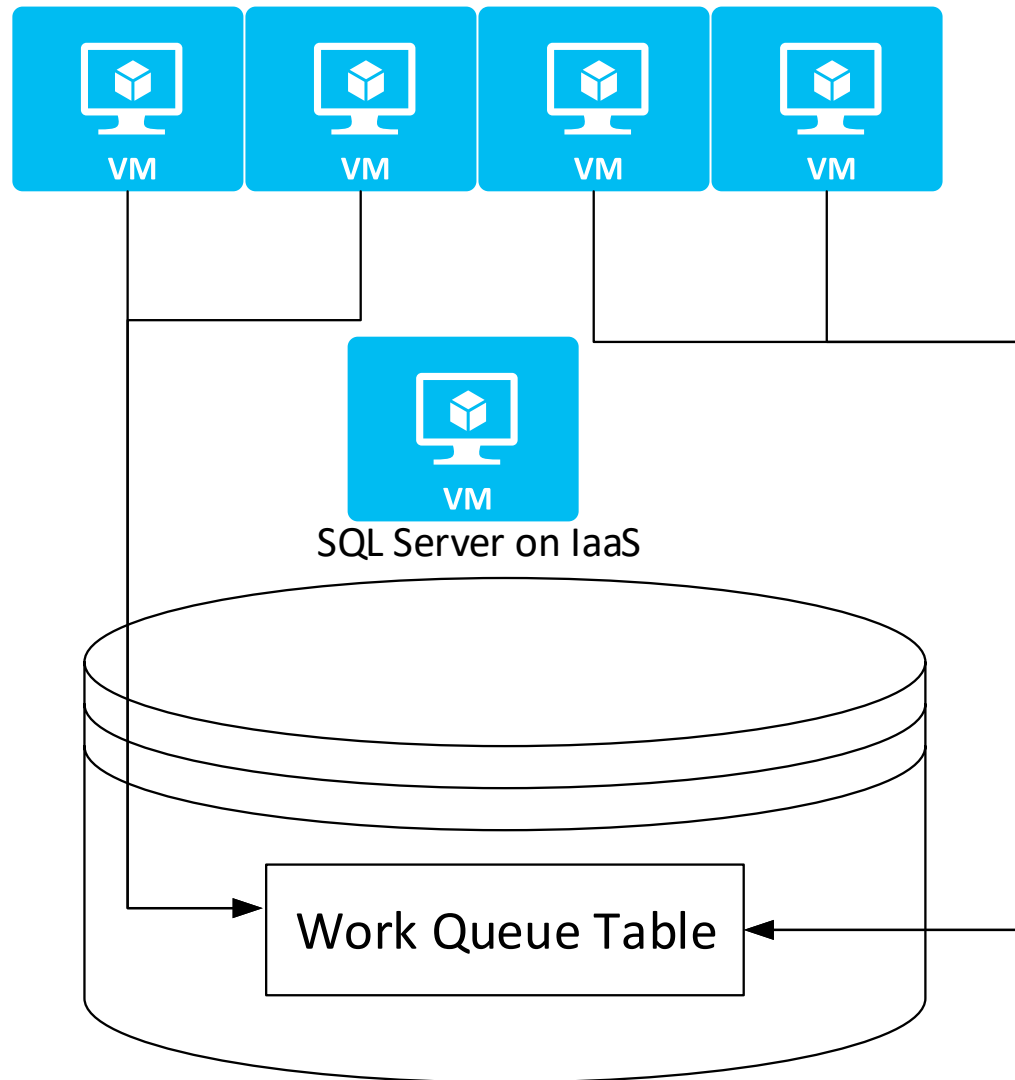
TalkTalk TV Challenges



TalkTalk TV Challenges - Duplicity




Monolithic Systems / Concurrency



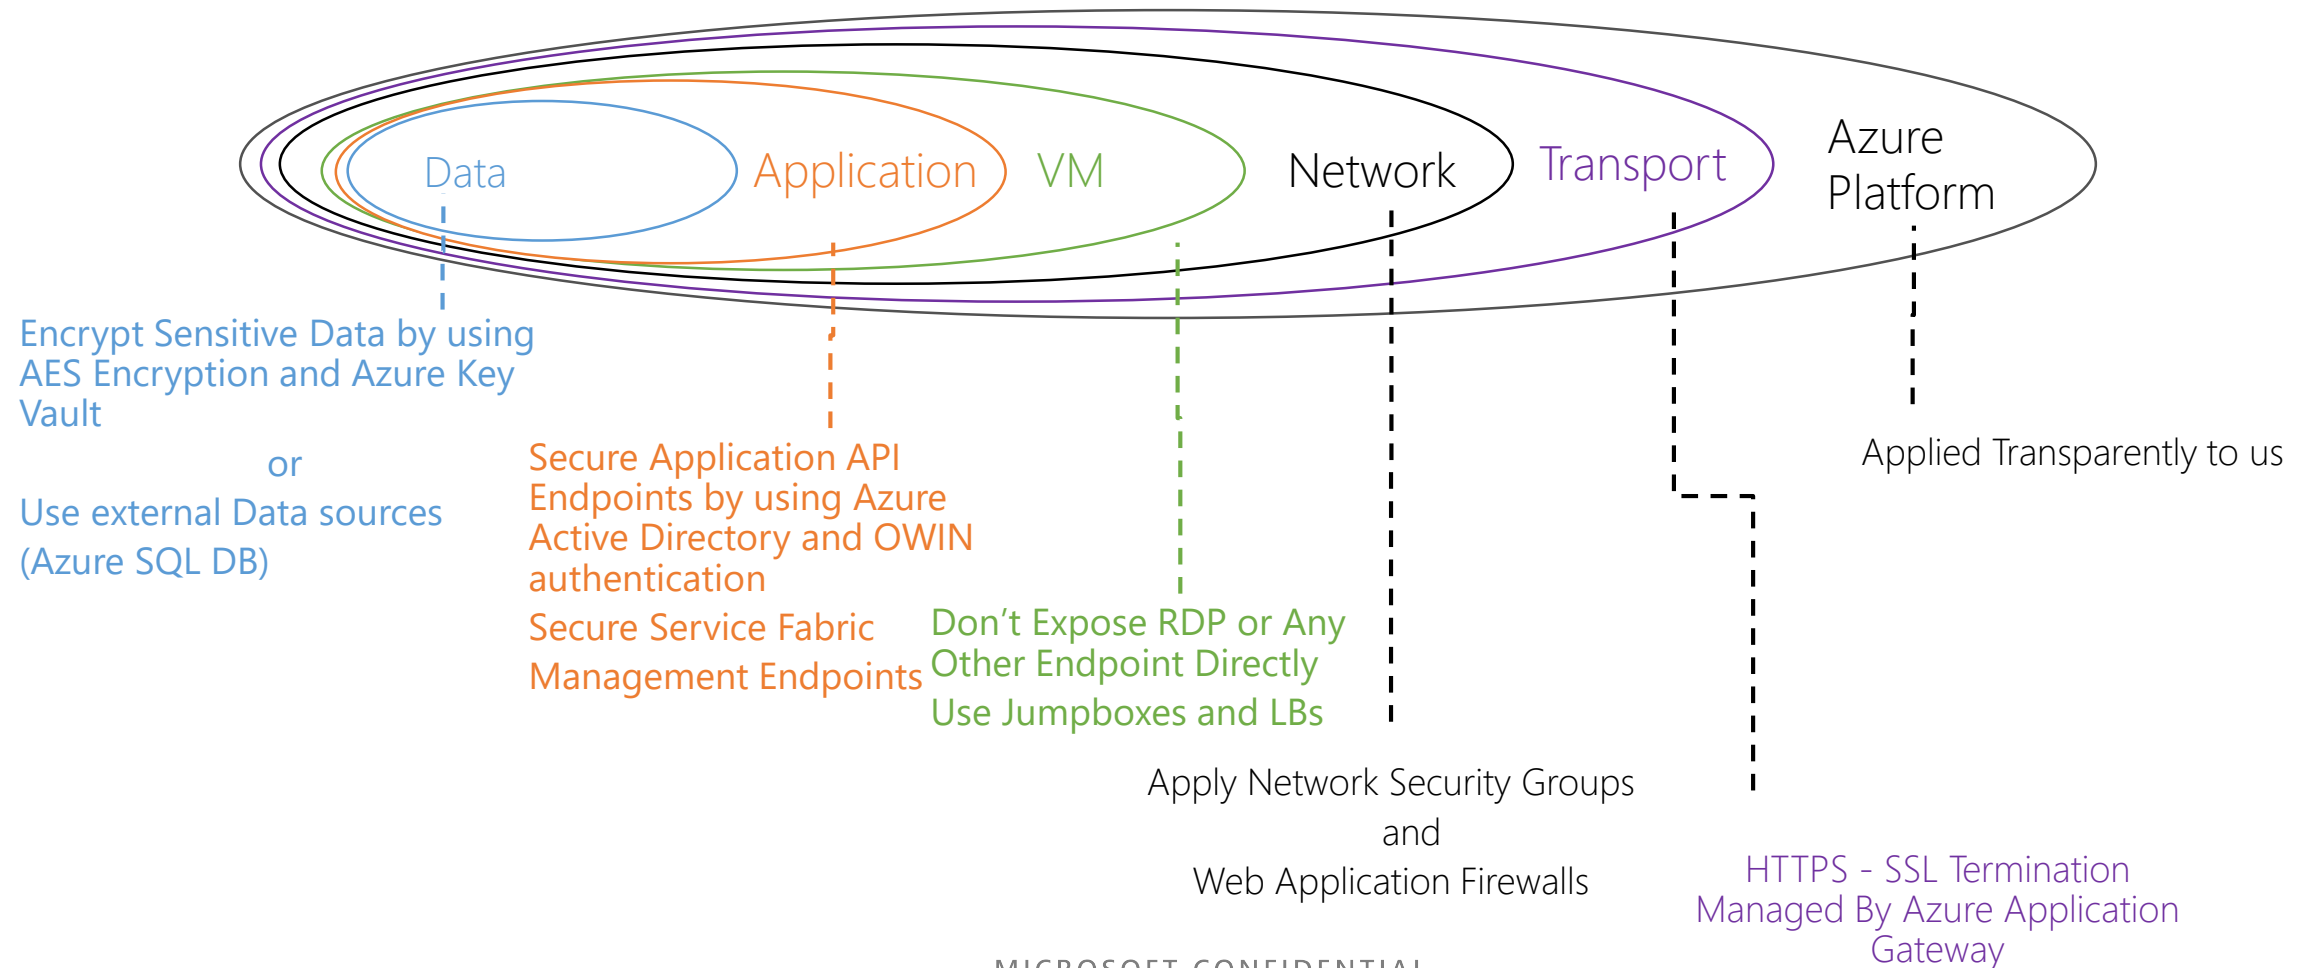
- During large batches of encoding cycles the DB suffered high contention
- Despite having an elastic platform (Azure Media Encoders), the ability to submit large encoding jobs was constrained by traditional design paradigms
- Too tightly coupled APIs / systems

Cultural change to development

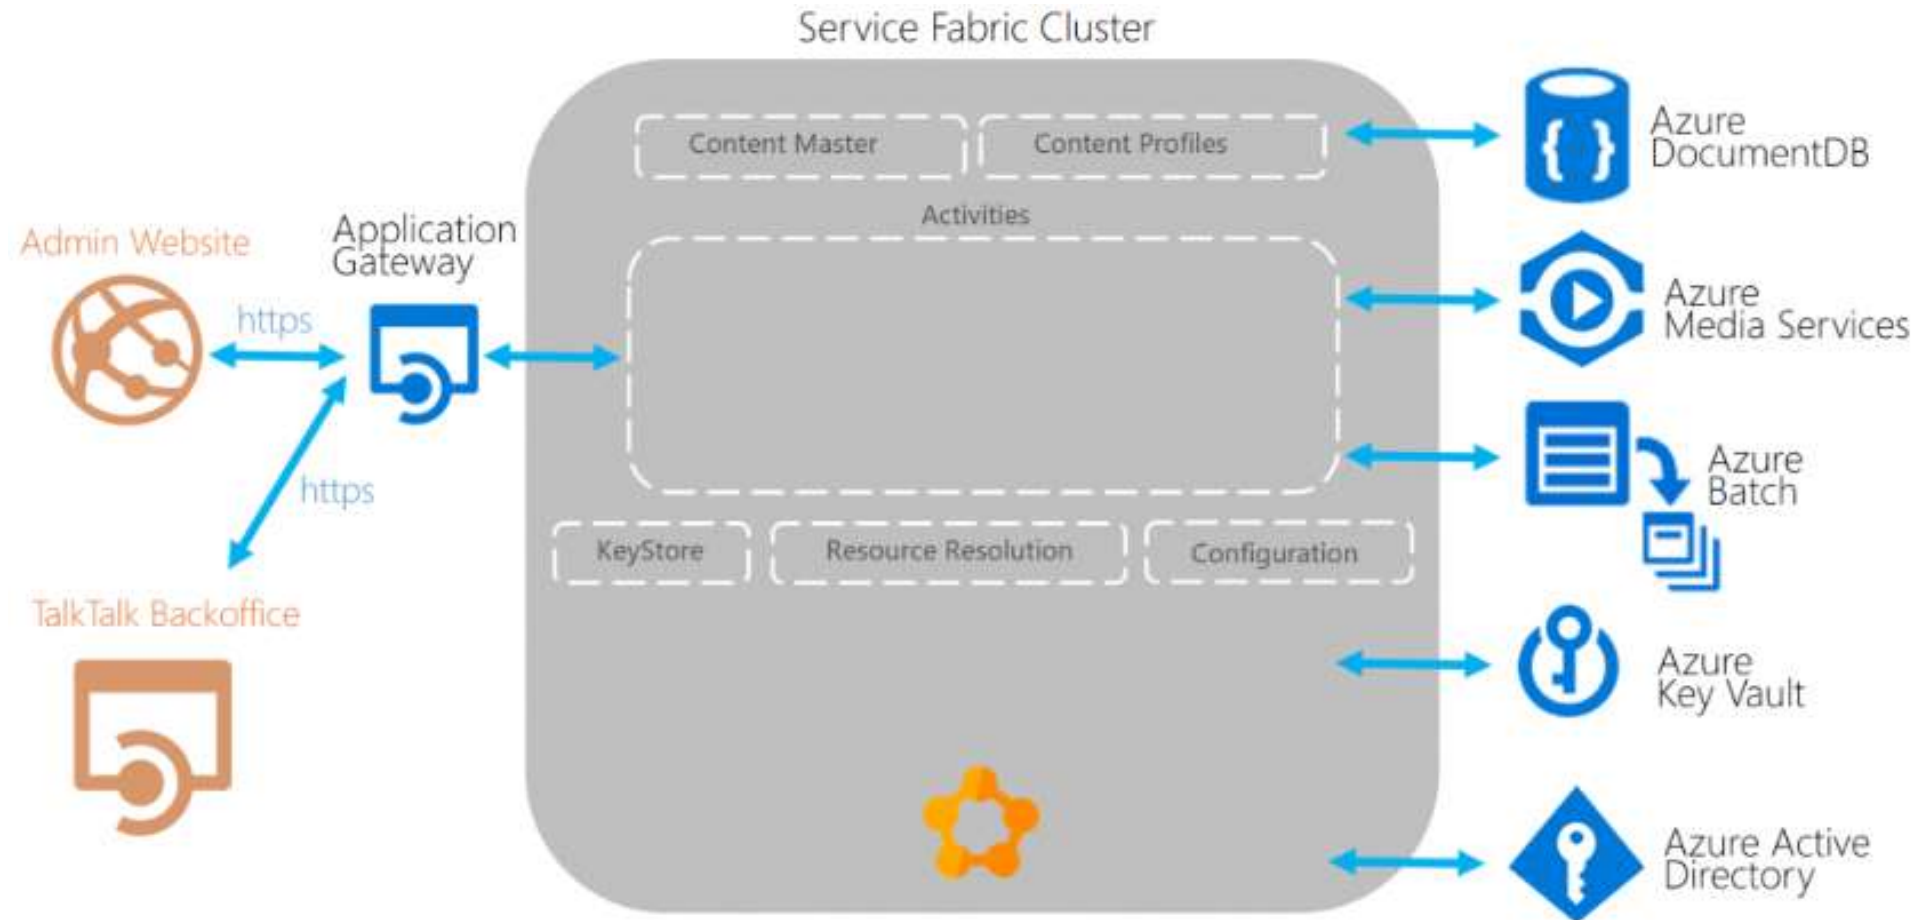


- Needed to adopt a more nimble and autonomous mentality
- Build small and cross functional teams
- Hydra : 3 engineers + Microsoft Support
- Developers + IT Ops – how to interact

Adopting an overall approach to security



Hydra on Azure Service Fabric – Conceptual Design



Hydra on Service Fabric – Authentication and Security



Azure AD - Applications

- Azure AD for user authentication to SF Apps
- Certificates/Secrets for cross-service authentication



Azure Key Vault

- Store Azure keys (Media Services, Batch, DocumentDB...)
- Configuration settings
- Certificates

Hydra on Azure Service Fabric – Media Jobs



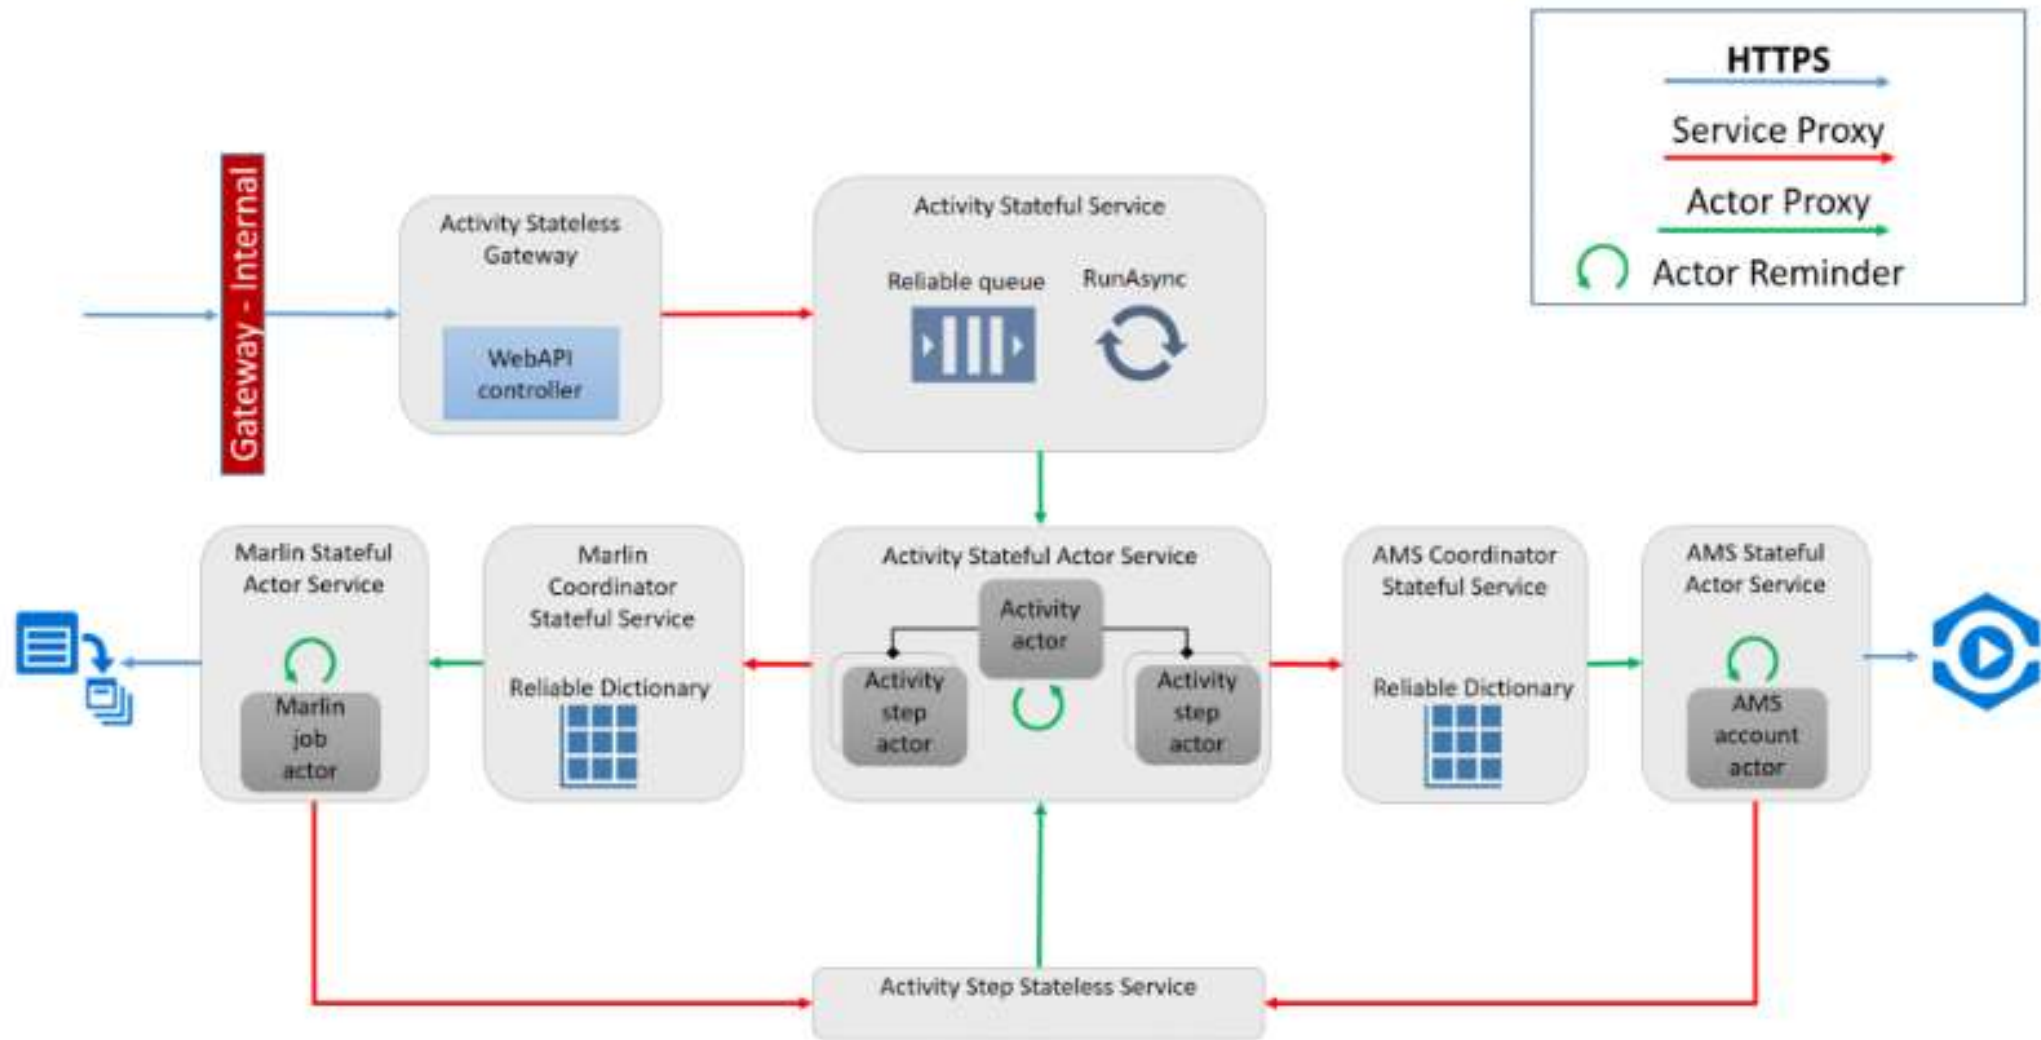
Media
Services



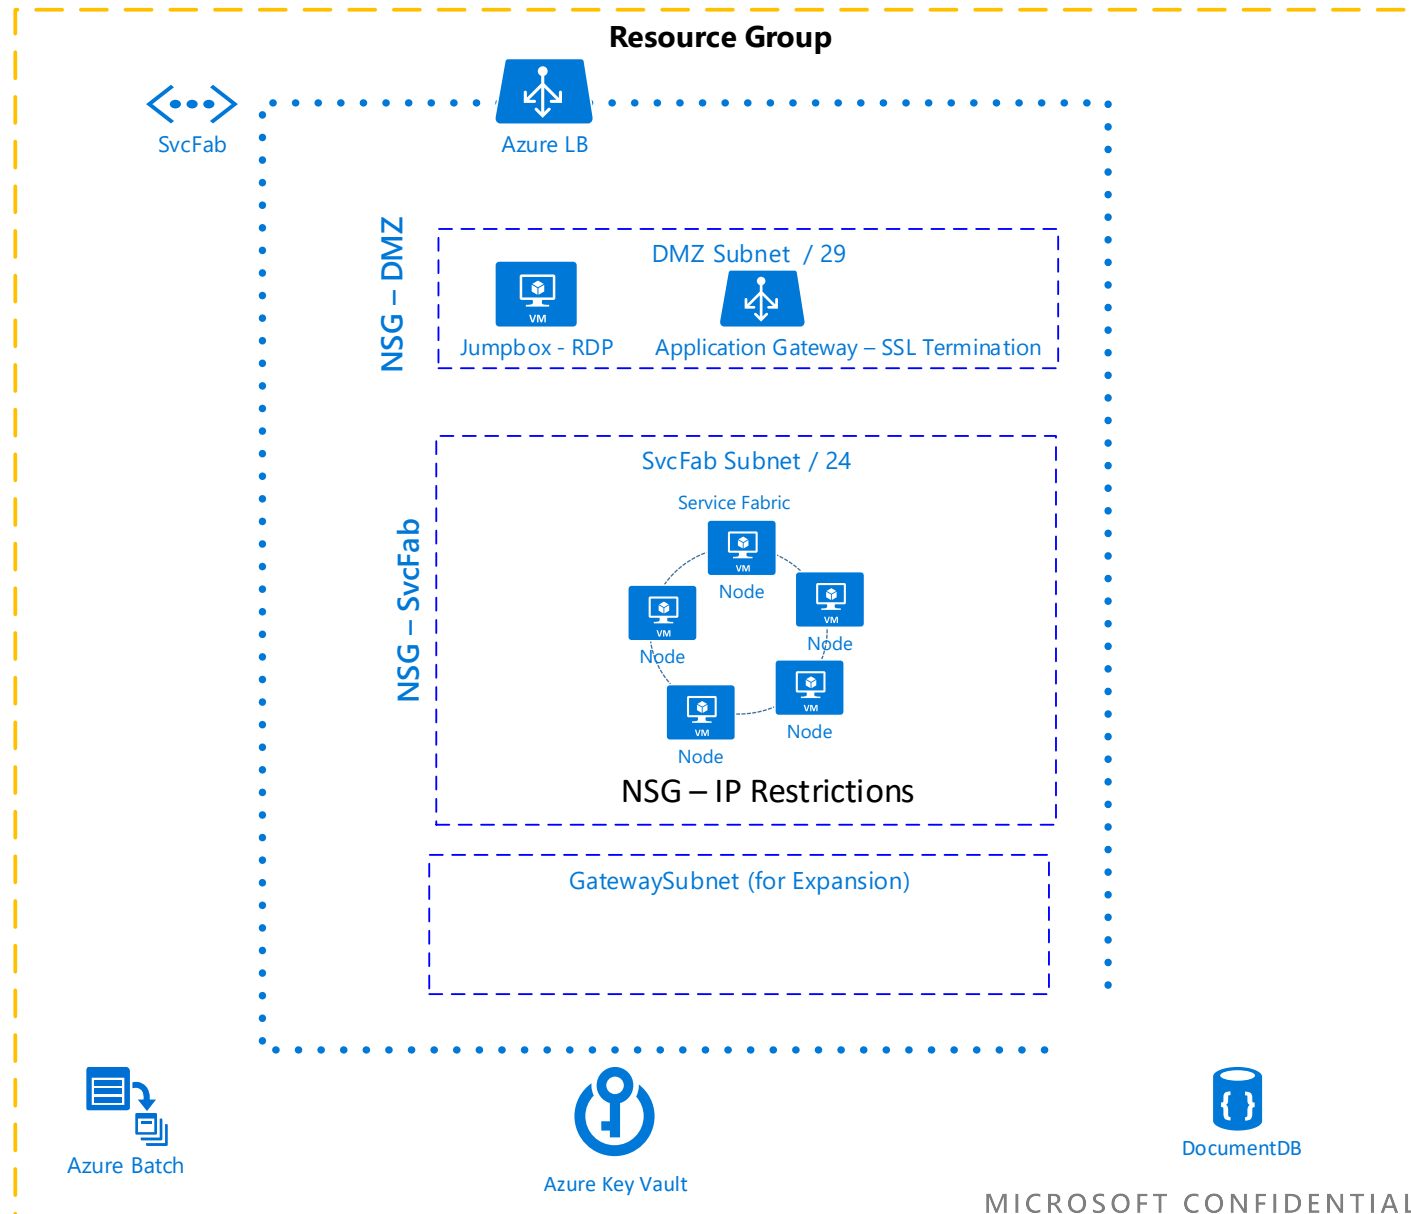
Azure Batch

- Azure encoders provide elastic encoding
 - Azure Streaming Services to deliver the content to IP based devices
-
- Extends Azure Media Services capabilities
 - Used for Marlin Encryption and long running jobs

Hydra on Azure Service Fabric – Detailed Services Design



Hydra on Azure Service Fabric – Infrastructure Design



- Use Azure resource manager templates for deployment
- Application of security policies :
 - SSL termination
 - Endpoints authentication
 - IP restrictions

Lessons Learned from TalkTalk

- Brown field development
 - You do not have to re-write the whole app to use microservices approach
- Pay attention to application and service composition
 - Loose coupling and high cohesion (contracts)
- Use the right programming model for each scenario
- Determine security and telemetry requirements up front
- Small team agility - The Perfect Combination
 - A Small Team (3 "+1") with a nimble engineering mentality
 - Right platform (Azure/ Service Fabric /Media Services/Batch)



Service Fabric - Use Case Scenarios and Examples

IoT on Service Fabric : Mesh Systems

Microsoft Services

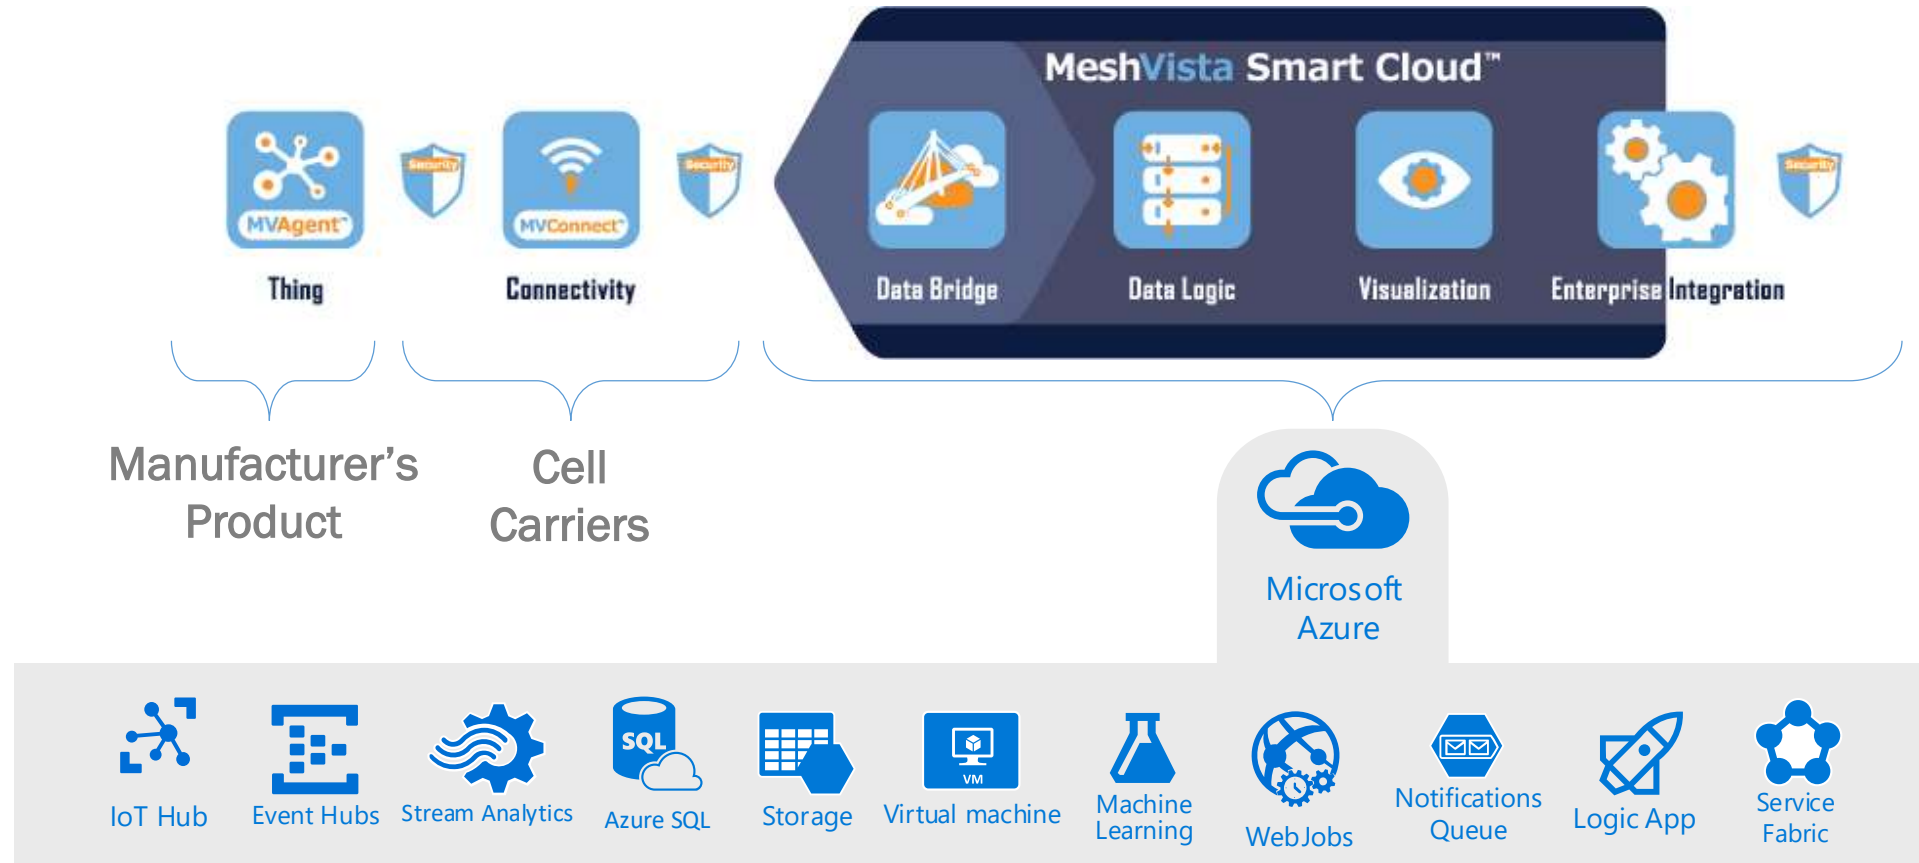


IoT On Azure Service Fabric

- Mesh Systems – NetLink
- Goals Of project
- Architecture
- Performance Metrics
- Lessons Learned

Mesh Systems

- IoT Software and Services
- Connect Devices To Our Platform in Azure



NETLINK

Smart Energy

Product Manufacturer: **WLS** LIGHTING SYSTEMS

Business Case: Reduce maintenance and energy costs by only running lights when and where necessary.

Device: Multi-channel Wireless Light Controls and Base Station

Azure : Data logging, remote scheduling, event notifications, energy usage analysis

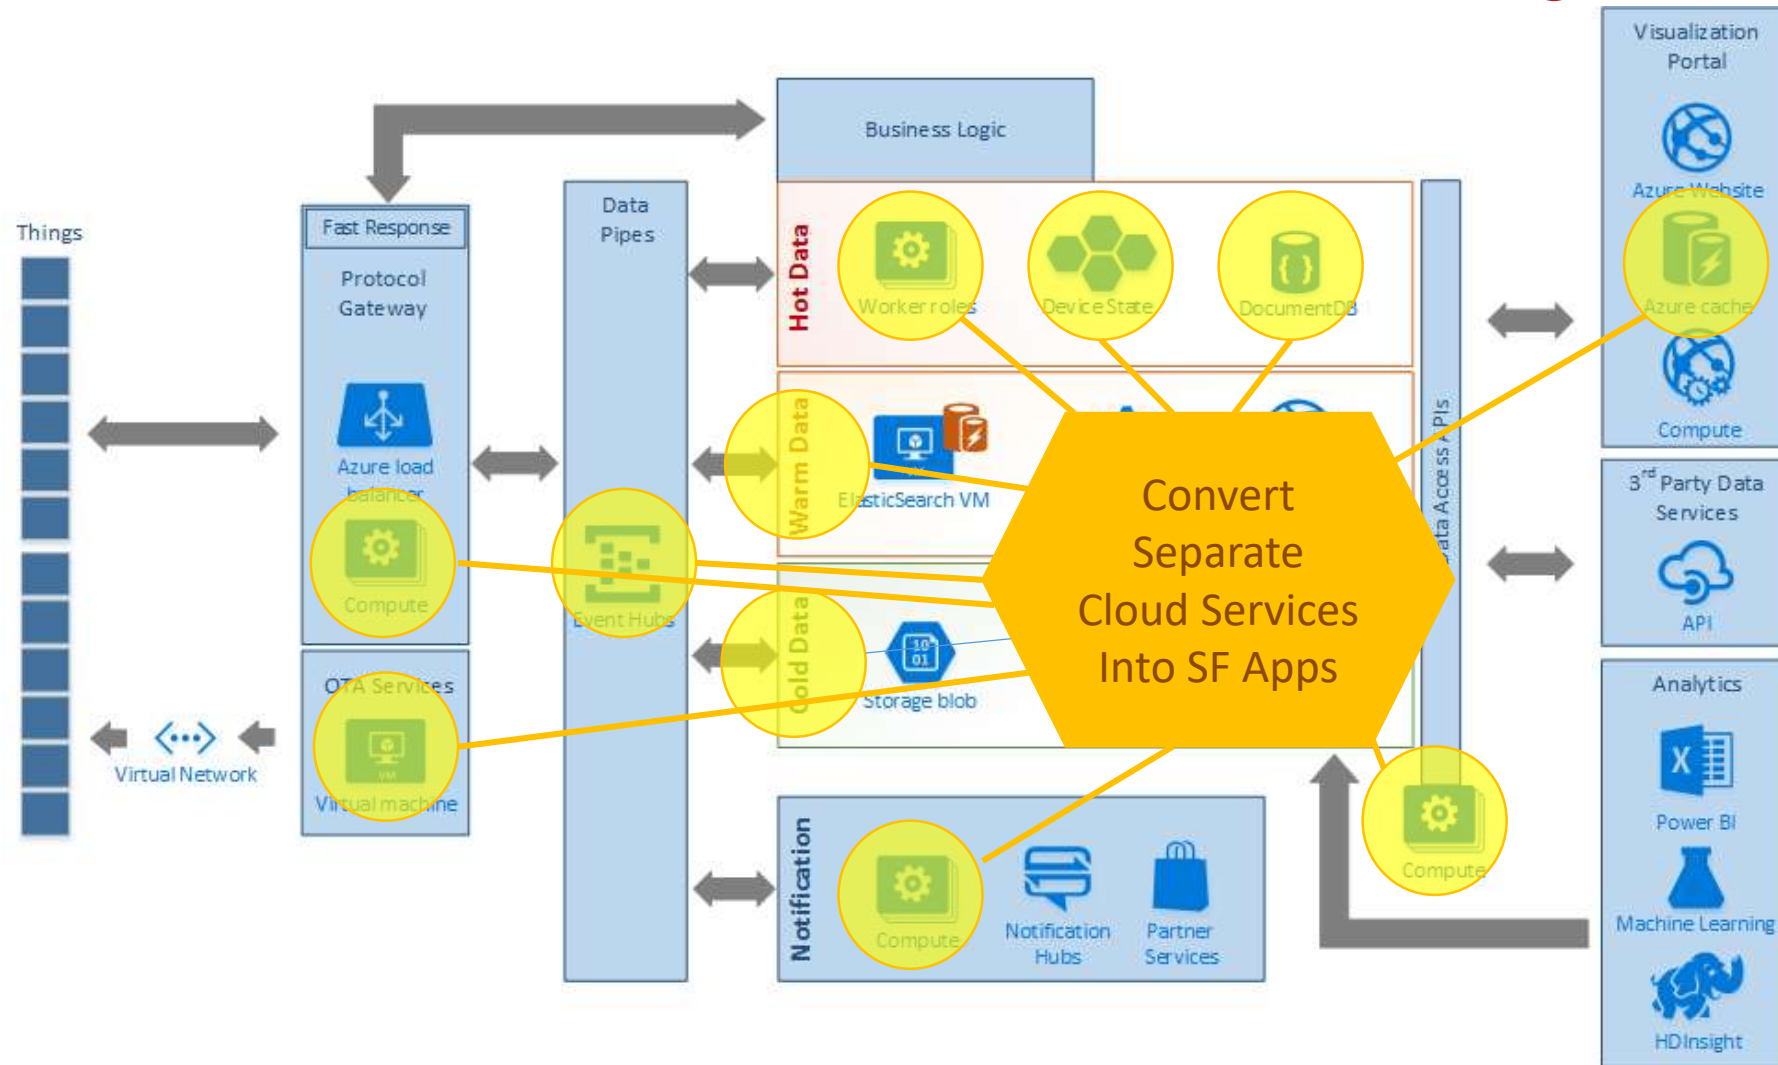
Key Comments:

- Clients see 12-19 month ROI
- Improved safety and quality of service



PaaS V1 Architecture


MESHVista®



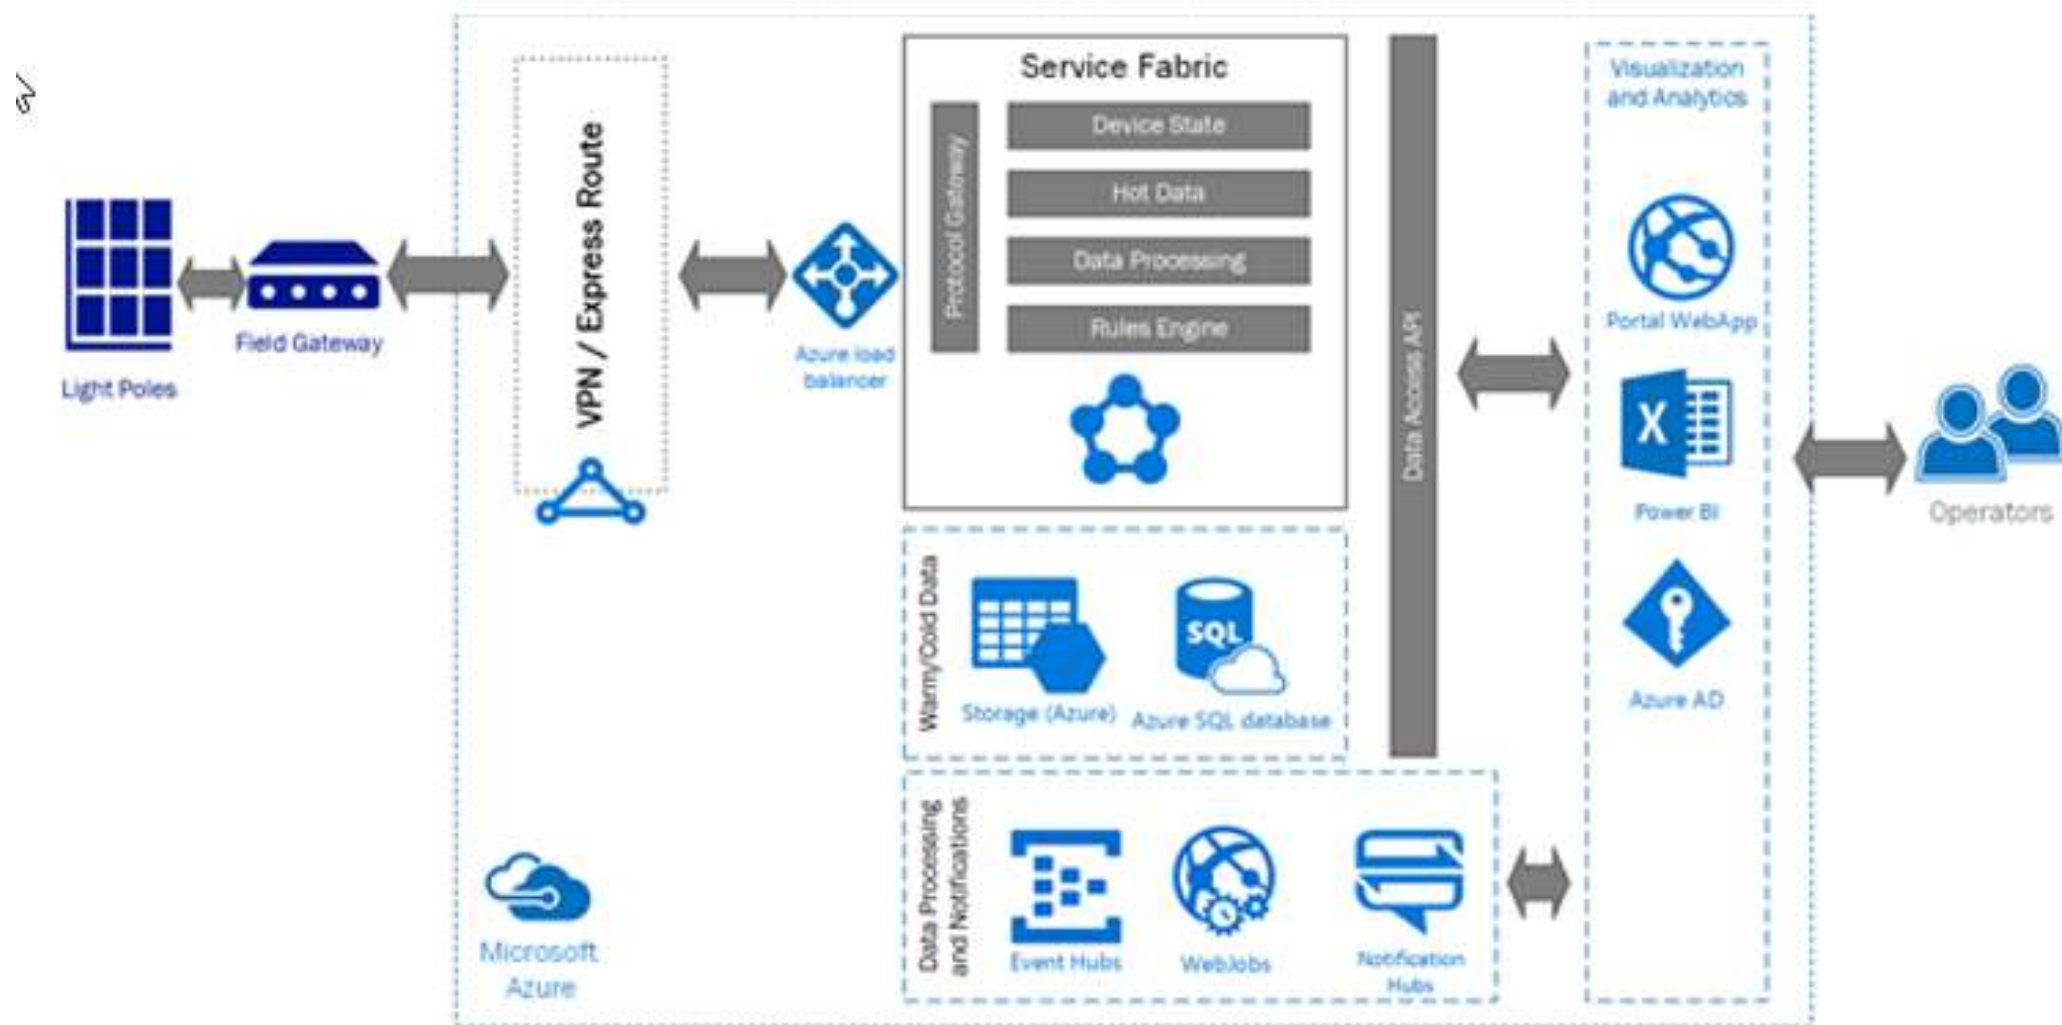
Goals of project

- Bare-Metal Migration to Azure
 - Aging Infrastructure needs to be replaced
 - No easy way to scale with old model
- Why Choose PaaS V2 vs V1?
 - Reduce cost and improve resilience
 - 3 Worker/Web Roles for each task to meet SLA
 - Role only running < 10% capacity
 - Device State / Actor Pattern – light fixtures represent Actors, not available in V1

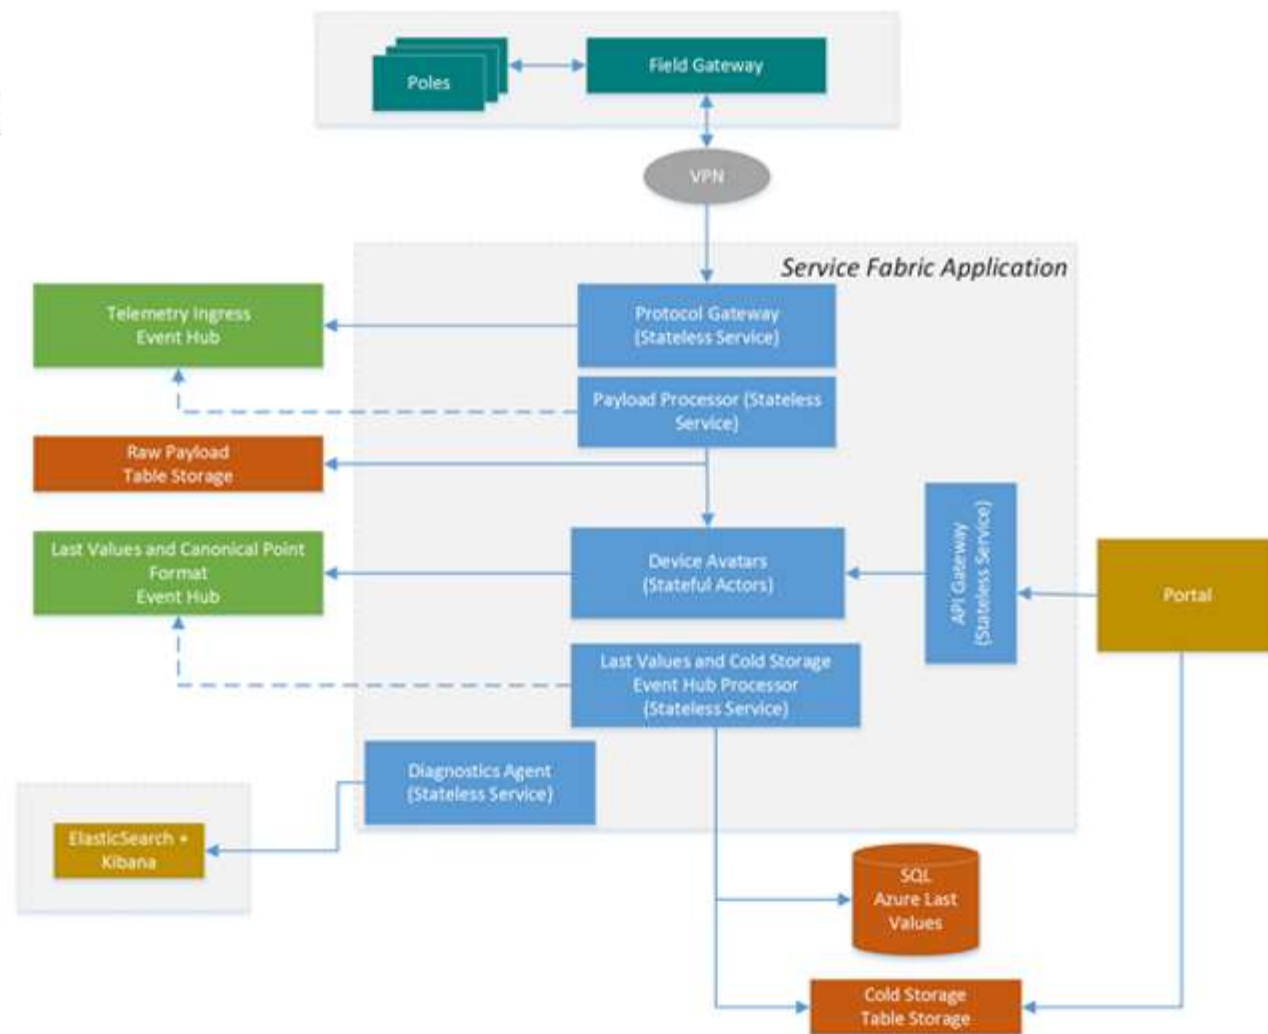
NetLink WAN – Data flow



MeshVista Hydra Architecture



Mesh Systems microservices



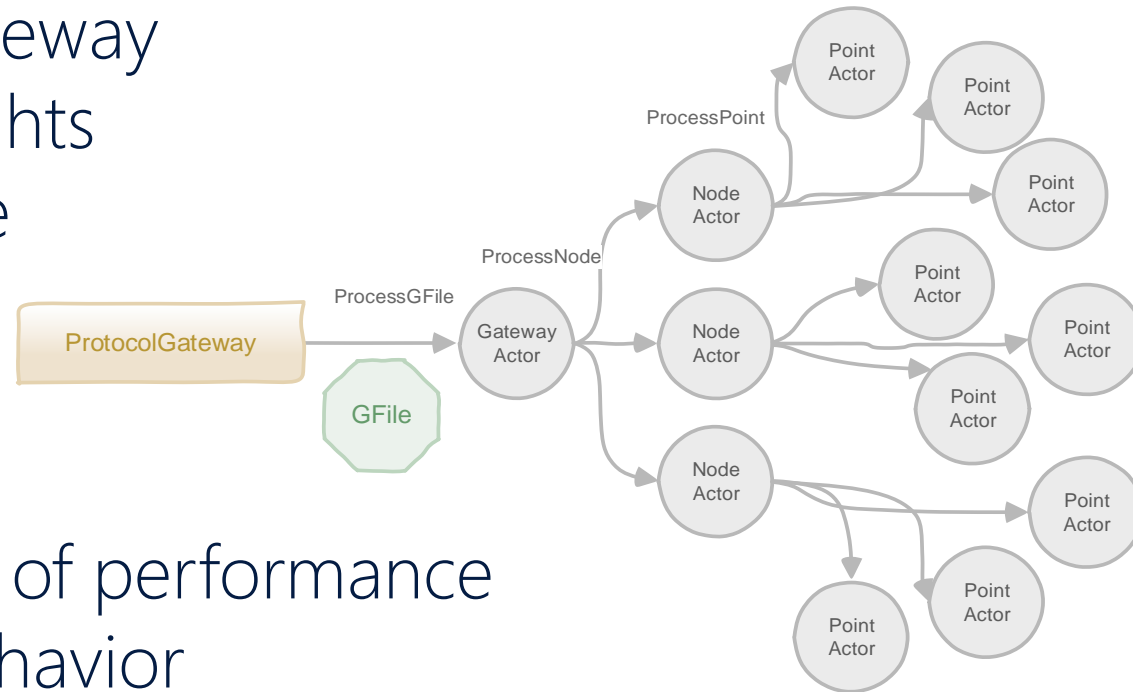
Device Actor Fanout

**Each incoming message evokes
a fan out of about ~200 actors**

Gateway: Field Gateway

Node: Pole with lights

Point: Sensor Value



- Simplicity of actor vs cost of performance
- Business rules of Point behavior (alarms, timers, aggregation rules, etc.)

Environment Applications

Protocol Gateway service – Parses the inbound device payloads

Device Avatars Actors service – Holds device state and runs device specific logic

Cold Storage Application – Persists device time series data to long term storage

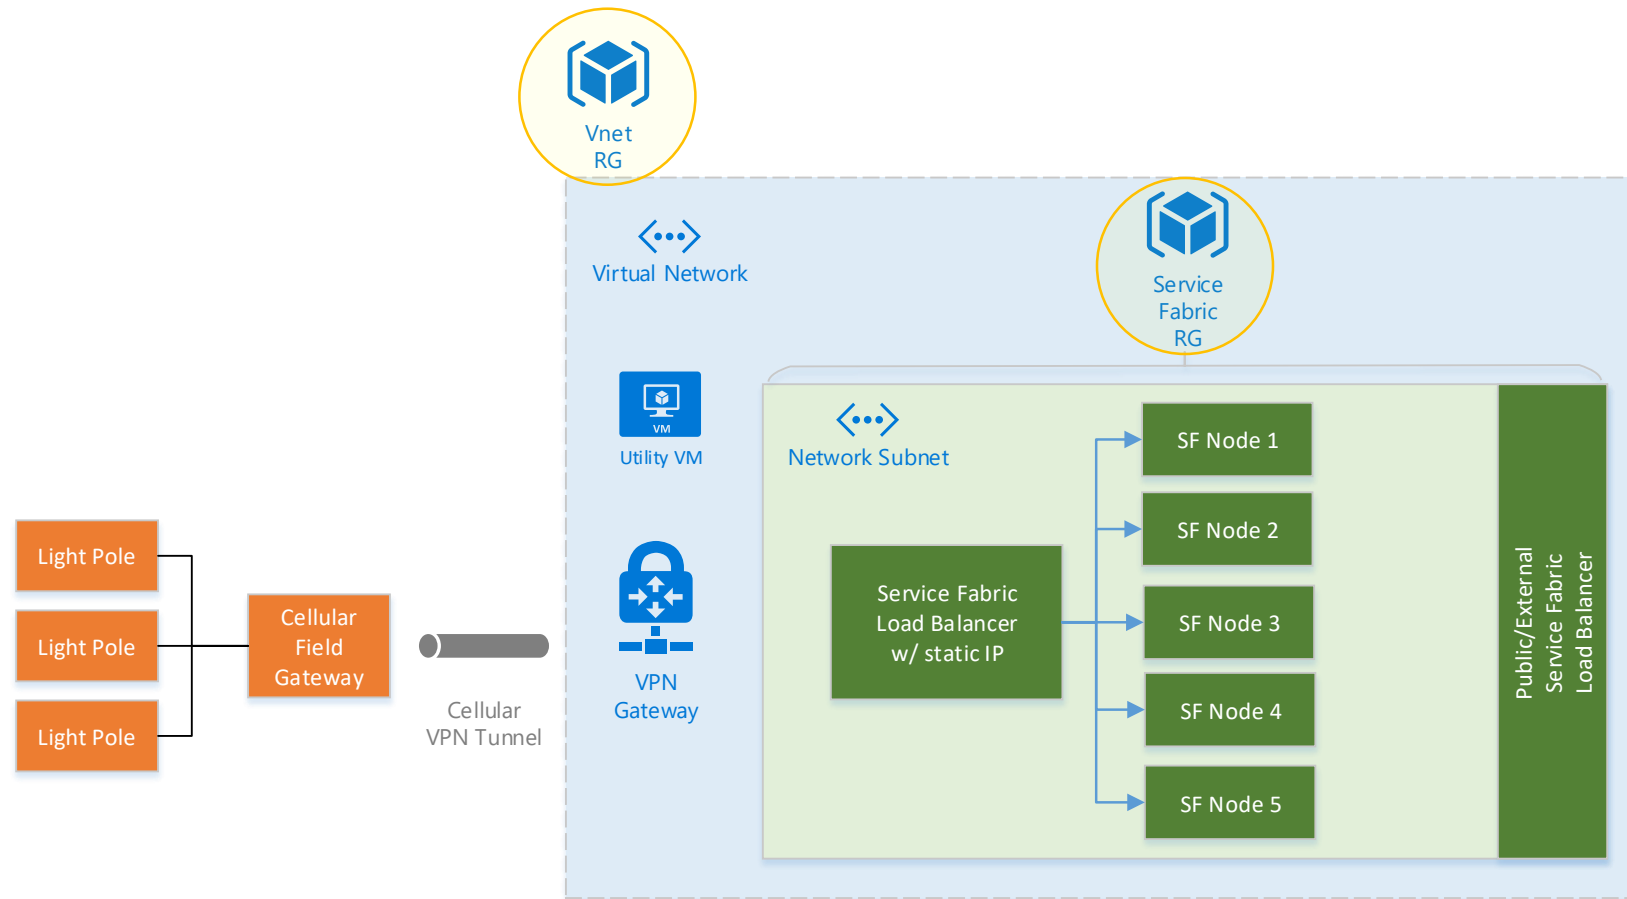
Last Values and Canonical Point Format Event hub service – Persists last value data to query-able storage

Raw Payload Processor service Persists raw, unprocessed device payloads to archive storage

Diagnostics Agent Service – ETW listener forwards app and health diagnostics to ElasticSearch

Elastic Search/Kibana Application – Automated Elastic Search index management

Two ARM Templates for deployment



Performance Metrics

- Class of machine A2 – A3 RAM (5 machines)
 - 300k actors were CPU intensive
 - RAM pegged on A2
- Ways to improve the process
 - Placement Constraints would help
 - Autoscaling Would Help Spikes
 - Azure Customer Advisory Team provide test result...

Details

- Test Scenario
 - 1,500 Gateways, 10 Nodes per Gateway, 20 Points per Node
 - Total: 300,000 actors
 - Gateway receives a GFile every three minutes
 - ~1,650 actors invoked per sec
 - Actor's State < 150 bytes
 - 10 minutes warm up (activations of all actors)
 - Measured latency of the full fan-out of a Gateway
- ServiceFabric test clusters
 - 2 test clusters: 3 D2 and 3 D3 machines
 - 3 services (GatewayActor, NodeActor, PointActor), each with 3 partitions and 3 replicas
 - KVS State provider, each invocation persists Actor's state

Lessons Learned

- Cohesive logging is critical and hard
 - Need big picture for distributed services
 - Semantic filtering is better than timeline
- Take time to pick VM class
- CPU/Memory 300k sensor actors
 - Fan out
- DB “Broker” is huge help
 - Minimize trips to the database

Lessons Learned (continued)

- Density/Cost - At small scale can have higher costs. At larger scale Service Fabric more cost effective
 - For MeshVista, 5 A3 machines cost more (approx.) 10-20% than existing cloud services solution.
- Be careful with actor programming model: Simplicity over performance
 - Careful of actor granularity. Don't go too granular or too coarse
 - Lots of message passing between actors will affect performance
 - Care of single threaded concurrency – calling “child” actor, blocks calls on “parent” actor. This can lead to longer fan-out times
 - Querying actors should be limited. Actors are great for isolated “avatars” but not good where data has to be aggregated or to build a view



Service Fabric - Use Case Scenarios and Examples

Illyriad Games

Microsoft Services





Independent games developer specializing in multiplayer games

Micro-studio with 5 core members of staff

Age of Ascent




- Born in the cloud – A new scale of gaming
- Space twitch combat, empire building, player driven economy
- Millions of players concurrently in single contiguous universe
- 50,000+ simultaneous players, real-time in the same battle
- 267 million network messages per second (23 Trillion per day)

“Age of Ascent is poised to become something special, one of those games that players will point to years down the road and say, *‘Yep, that was the one that changed everything.’*” - MMORPG.com

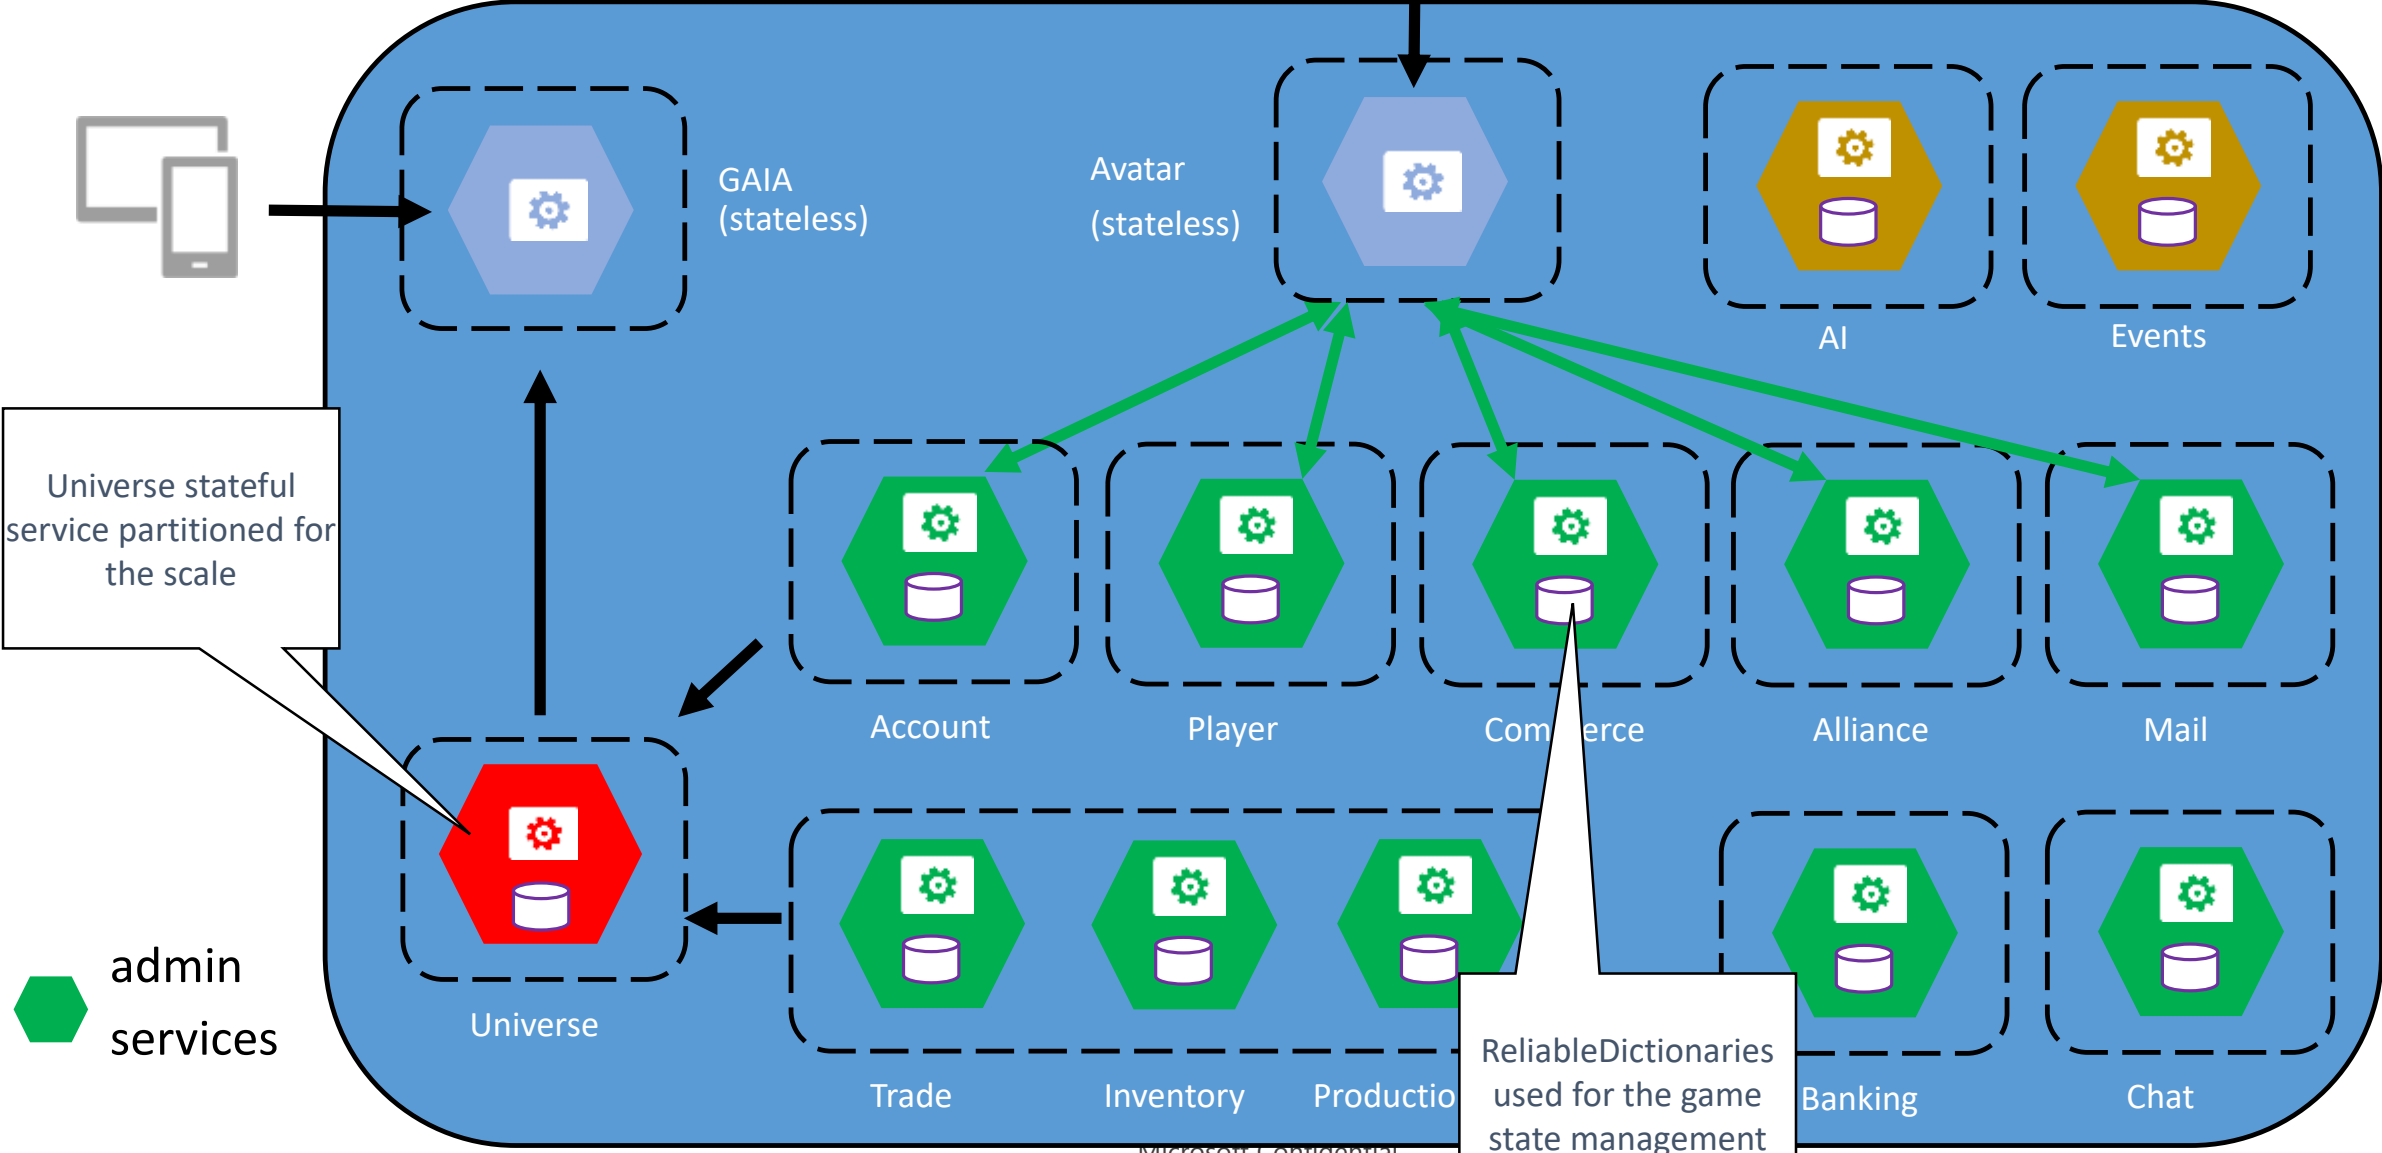
Age of Ascent Technology

- Use Microsoft Azure for hyperscale
- ASP.NET Core x64
 - Platform Agnostic – HTML5 + JS + WebGL Client, Azure Backend
 - Anywhere, Anytime – Cross device
- Service Fabric – game state management
- Listen to Hanselminutes podcast - Inside Age of Ascent with Ben Adams
 - <http://hanselminutes.com/509/inside-age-of-ascent-with-ben-adams>

Age of Ascent Architecture



Service Fabric Cluster



Service Fabric Benefits for Age of Ascent

- Deployment and upgrade agility
- Cost and efficient utilization due to high density
 - Illyriad estimated they would need ~700 cloud service machines to run Age of Ascent. On Service Fabric can achieve the same with < 100
- Elasticity scale in and out - manage costs
 - Dynamic Scaling – Spatial unfolding and refolding with scale sets
- Cluster management and orchestration
 - Dynamic Load Distribution – Service Fabric balancing of hot and cold nodes
- Always available stateful microservices

Key Service Fabric Features

- Actors = In memory device avatars
- Lifecycle management
 - Rolling app independent updates
- Application specific horizontal scalability
 - Scale apps independent of compute (Density)
- Reliable collections – cache in place
 - Database load/connection throttling
- Lightweight applications by concern

In Review: Takeaways

- Always have a clear and concise understanding of what you mean by scale and agility for your particular scenarios
- Service Fabric is a platform for building applications with a microservices approach
- See other case studies at <http://blogs.msdn.com/b/azureservicefabric/>

Lab: Using Service Fabric to build a customer order application



